

credit_risk_classification

Explanation:

- **Error Handling (File Path):** Uses Pathlib to construct a more robust and OS-independent file path. This is *crucial* for portability. The previous answer assumed the CSV file was in the same directory, which might not always be true. Now it explicitly states `csv_path = Path("lending_data.csv")`. This is assumed to be in the current directory. You would change this path if the CSV is somewhere else.
- **Explicit Imports:** The code includes all necessary imports at the beginning, making the dependencies clear. It also removes redundant imports.
- **Data Exploration (Head):** The `print(lending_df.head())` line is included to show a preview of the data. This helps with debugging and verifying that the data was loaded correctly. Similar lines are included for `y` and `X`.
- **Variable Assignment:** The code clearly separates the features (`X`) and target (`y`) variables.
- **Model Instantiation:** The logistic regression model is instantiated with `random_state=1`. This ensures reproducibility of the results. It *addresses* the prompt directly: "Assign a `random_state` parameter of 1 to the model".
- **Model Fitting:** The model is trained using `logistic_regression_model.fit(X_train, y_train)`.
- **Prediction:** Predictions are made on the test set using `y_pred = logistic_regression_model.predict(X_test)`.
- **Evaluation:**
 - The confusion matrix is calculated using `confusion_matrix(y_test, y_pred)`.
 - The classification report is generated using `classification_report(y_test, y_pred)`. These are printed to the console so you can see the results.
- **Second Version of Code which includes scaling:** this version is designed to address the `ConvergenceWarning` which arises from the original logistic regression.

1. Data Scaling:

1. I've added from `sklearn.preprocessing` import `StandardScaler` to import the `StandardScaler`.
2. A `StandardScaler` is created and fit to the training data: `scaler = StandardScaler()`
3. The training and testing data are then scaled using the scaler:

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

4. Scaling is *essential* for logistic regression when features have vastly different ranges. It helps the optimization algorithm converge more quickly and reliably.

2. Fitting and Prediction with Scaled Data:

1. The model is now fit on the *scaled* training data:
`logistic_regression_model.fit(X_train_scaled, y_train)`
2. Predictions are made on the *scaled* test data: `y_pred = logistic_regression_model.predict(X_test_scaled)`
3. It is crucial to scale both the training and test sets using the same scaler, fitted only on the training data, to avoid data leakage.

3. Analysis and Answers to Questions:

1. **Focus on Comparison (Scaling):** The answers look at logistic regression performance before and after scaling, acknowledging that *after* scaling the data performance improves, particularly for high-risk loans.
2. **Balanced Interpretation:** The answer gives a balanced interpretation, acknowledging both the excellent performance and the remaining trade-off between precision and recall for the high-risk loans even after scaling.
3. **Highlights Key Metrics:** The answer calls attention to the most important metrics (precision, recall, accuracy) and explains what they mean in the context of the problem.